

Walmart Java Guild Presents:

# GraphQL



By John McCaulley



Quick Intro – About me!



## Topics

- **What is GraphQL?**
  - Basics
  - Interfaces/Unions
  - Request/Response Structure
  - GraphQL
  - Execution
  - Resolvers
- **Data Loaders and N + 1**
- **Java Frameworks**
- **Schema Design**
- **Aliasing vs Batching**
- **Versioning (or not)**
- **Directives**
- **Errors**
- **Security concerns**

Don't forget to pace based on how familiar everyone is.

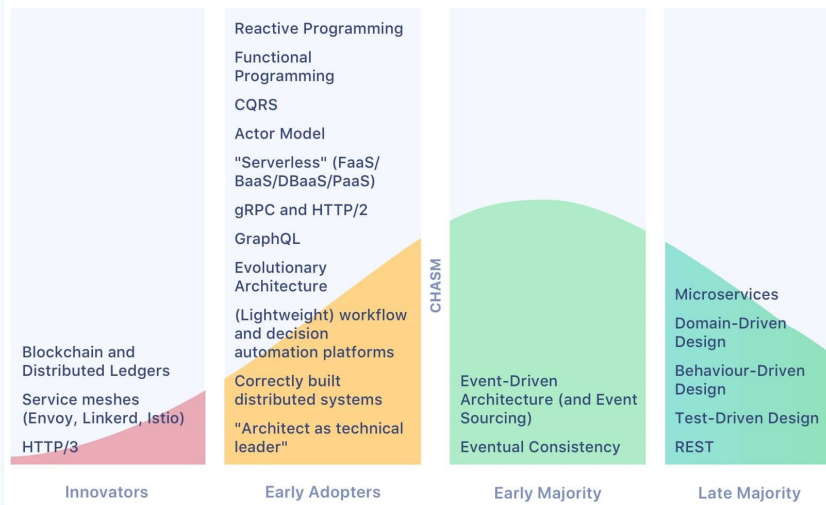
Going to cover basics with the goal of giving everyone a good understanding of why data loaders are necessary and why we need to deal with the problems they create in java/schema design.



## Software Development Architecture and Design 2019 Q1 Graph

<http://infoq.link/architecture-trends-2019>

InfoQ




# What is GraphQL?

“GraphQL is a query language for APIs and a runtime for fulfilling those queries with your existing data. GraphQL provides a complete and understandable description of the data in your API, gives clients the power to ask for exactly what they need and nothing more, makes it easier to evolve APIs over time, and enables powerful developer tools.”

First ask the crowd

⚡ Cool... but what does that mean? ⚡

⚡ And why “Graph”QL? ⚡



Like REST, GraphQL sits on top of http and defines an “API design architecture”. One important part is "gives clients the power to ask for exactly what they need". This is in direct contrast to REST. Instead of resources, everything is a graph of relationships, aka why the graph. The other important part is the Runtime - I would argue is much more structured than REST and it has to be because GET, POST, UPDATE, DELETE are not enough.



Really, three major parts:

- Type System
- GraphQL schema language (SDL)
- Execution Semantics

Type systems all start with the special root, service specific set of possible relations to query. Defines why returned data can look like and allows validation of queries

SDL, schema definition language, programming language agnostic way to talk about type systems

Execution semantics - GraphQL documents, differences between query and mutation, response map format (data, errors, extensions)

 Some basics:

```
type Character {  
  name: String!  
  appearsIn: [Episode!]!  
  height(unit: LengthUnit = METER): Float  
}
```

Cover the basics here as much as possible:

Everything is just a graphql object type - meaning it is a type with fields

Everything starts with the special root object and ends in scalars

Scalar types: String, Int, Float, Boolean, ID - same as a string but signifies may not be human readable

Enums are special type of scalar - language dependent in implementation.

Modifiers - Lists and Nullity (return type even has nullity)

2 special types -> query/mutation. Query is entry point and at least one required, mutation optional

Fields can have parameters

Parameters can have defaults



## Interfaces, Unions, Fragments:

```
interface Character {  
  id: ID!  
  name: String!  
  friends: [Character]  
  appearsIn: [Episode]!  
}
```

```
type Human implements Character {  
  id: ID!  
  name: String!  
  friends: [Character]  
  appearsIn: [Episode]!  
  starships: [Starship]  
  totalCredits: Int  
}  
  
type Droid implements Character {  
  id: ID!  
  name: String!  
  friends: [Character]  
  appearsIn: [Episode]!  
  primaryFunction: String  
}
```

```
{  
  leftComparison: hero(episode: EMPIRE) {  
    ...comparisonFields  
  }  
  rightComparison: hero(episode: JEDI) {  
    ...comparisonFields  
  }  
}  
  
fragment comparisonFields on Character {  
  name  
  appearsIn  
  friends {  
    name  
  }  
}
```

```
union SearchResult = Human | Droid | Starship
```

Interfaces are pretty self explanatory, unions are groupings that have nothing in common, fragments can save space or be inline - they are conditional  
We haven't used either and from what I've read interfaces are the only ones people find really useful.

## Request structure (in JSON)

```
{"query" : String, "operationName" : String, "variables" : Map<String, String>}
```


## Response Structure (again, JSON)

```
{ "data" : Map<String, String>, "errors" : [...], "extensions" : Map<String, String>}
```

Interesting points - this is designed with lightweight browsers in mind. Query isn't really a query, it is a static set of operations, op name lets you choose which one - avoiding string parsing etc on the client

Query, Mutate are by convention similar to get, post in REST. also have subscriptions  
Subscriptions are not durable like kafka etc might be, but more for things like news feeds. Java implementations are socket based.

# GraphiQL and Introspection - SIRO!



Show siro's graphiql - maybe demonstrate variables etc...

## Execution

- Parse > Validate > Execute
- Everything is a type, which is some hierarchy from the root
  - So really, everything is just a field of the enclosing type
- Each field is associated with a function to obtain the type - a *resolver*
- Once a resolver completes, arguments and the results (usually called source) supplied to next level in the hierarchy

Validation can occur at any point with the type system - ie, after introspection a client can validate their own queries before sending

We are not getting into details of the "Runtime". Resolving happens down until leaf nodes or scalars.

Resolvers can be anything - arbitrarily complex, many though are as simple as a property - these are usually supplied by the implementation by default - as we will see with java-graphql.

This is in large contrast to REST, you could view a rest endpoint as a single resolver.

Can anyone see how this is beneficial? How might it suck?

You must have a resolver per field because you don't know what will be requested

Resolving sounds simple right? Any obvious problems?

A naive resolver example:

```
query {  
  authors {  
    name  
    address {  
      country  
    }  
  }  
}
```

1st trip to the db gets the authors

2nd trip gets the address for each author  
=  $n + 1$  authors

For this example assume the author backing object has just a name/id field and the address resolver must make another db call.

## Data Loaders to the rescue!

- Developed by Facebook
- Asynchronous - revolves around futures/promises
- Utilizes a cache
- Supports batching
- Created specifically for javascript
- `Dataloader<T, S>.load(T t)`

Important to note that you want the batching/caching to be request scoped  
Can anyone guess why it starting as a javascript library might impact us?

## Current Java Frameworks

- GraphQL-Java: Core framework. Port of a js reference implementation.
- GraphQL-java-tools: A schema-first tool for graphql-java inspired by graphql-tools for JS
- GraphQL-java-servlet: Servlet endpoint for GraphQL Java
- GraphQL-spring-boot: GraphQL and GraphiQL Spring Framework Boot Starters
- Java-dataloader: A Java 8 port of Facebook DataLoader

These are the ones we are currently using and have experience with  
Notice everything is a port....

## Java GraphQL

- Schemas can be in code or in SDL
- Creating an executable schema requires a Runtime wiring (how to resolve each field) along with a definition of the types to resolve
- If resolvers not present in wiring, default is a property resolver
- No requirement for wirings - if nothing can be/is wired up null is returned
- Minimally supports Facebook data loader port.

## Java GraphQL Tools

- Automates data fetcher creation
- Enables IOC frameworks
- Hierarchy of resolvers: Object property > Object method > Resolver method
- Runtime type checking
- Uses Reflection

After reviewing this step into code examples leading up to the data loader dispatch problem

Not going to spend a lot of time covering the other frameworks libraries - they are kinda self explanatory

The tools project does classpath scanning to create data fetchers based on schema files (.graphqls)

Relate GOG example to BI and XREF



## Dataloader Pains

- Javascript version uses the event loop
- Java GraphQL's default is through instrumentation
  - Instrumentation build into graphql java aop like
- Dispatching is about competing concerns - being async vs lazy vs sequential relationships.
- Problem even more complex when request batching
- Is this a problem of the framework or a modeling problem?

It is important to note - JS does not have any of these problems

There is no current solution in java - this could be a good place for contributions

Further explain request batching - and how this isn't an issue with REST

Is the problem the framework is a port? Is it data modeling? Leading to the first "controversial" topic...



## Schema driven dataloading VS Dataloading driven schema?

Big point of contention, SIRO did the schema first. What is business logic within the API orchestration layer? Why would writing the db first be an anti pattern and writing the schema first not be? Best argument against so far is how to handle if you change the loading pattern but how often do you change data from one place to another, that changes the loading pattern (hint we never have)? In that case why wouldn't you simply treat the new one as a new field and the old one as deprecated?

## Aliasing vs Batching

Alias

```
{heroOne:getHero(name:"Thor") {powers} heroTwo:getHero(name:"Star-lord") {powers}}
```

Batch

```
[{getHero(name:"Thor") {powers}}, {getHero(name:"Star-lord") {powers}}]
```

If we didn't mention aliasing earlier explain here.

Aliasing is nice, but it destroys caching - can no longer cache document parsing, and can your cdn handle it?

You no longer get metrics the way you think you would

The client now has to maintain/create names

This sounds good right? Well not in Java - highlight how the instrumentation wasn't allowing batching to really happen.

Is the answer to add a new method that accepts a list instead (like REST would require)?

## Versioning


- Why do we use versioning in the first place?
- GraphQL - don't do it
- Deprecate using @Deprecated (provided by apollo and graphql java)
- Use tracing to determine if removal is safe

We version APIs to gain control over the shape of data  
GraphQL is highly opinionated about not versioning, you shouldn't but you can if you really want  
Additions are non-breaking.  
Let this lead into directives

## Directives

- @include(if: Boolean)
- @skip(if: Boolean)
- Others provided by frameworks (deprecated, cacheControl, deferred etc)
- Mostly useless for Java/Walmart use case

They seem powerful. Explain experience with chainRestriction and cacheControl and how they don't really work with ccm etc.

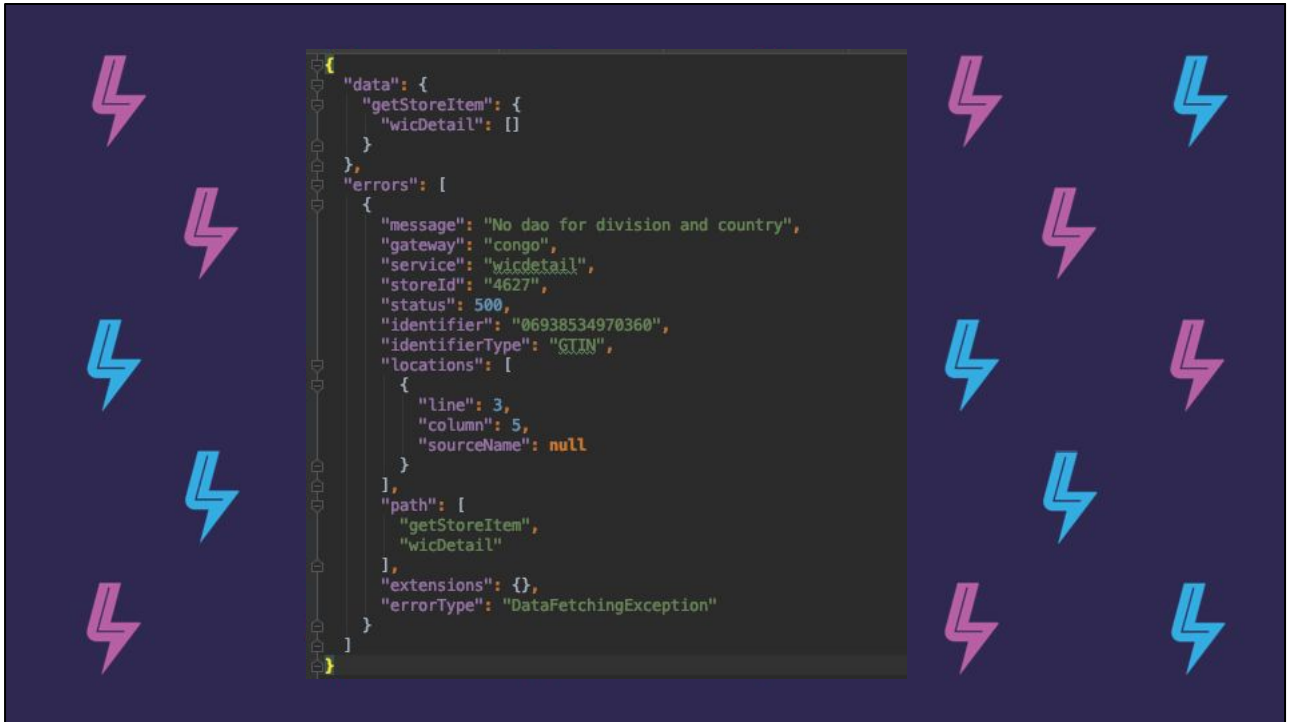


GraphQL errors are a huge growing pain, 200 only really means request got to the server and parsed.

Error handling is on the top of everyone's "5 things I learned doing graphql"

Have to keep in mind - graphql supports both errors and data at the same time

If you want the client to use it, think about making it a property or flag allowing general errors to use the built in error mechanisms



This is one of our old errors that we used to send back to clients. We are simplifying our error responses, but we still do not have a good way to communicate what errors are fatal etc.

## Security

- Enterprise data access not really a GraphQL use case
- Authentication on a per field level, how would it scale?
- How do you control what parts are externalized?
- Enter persisted queries

GraphQL is really designed with browsers in mind, you can authenticate but do things like what fields to externalize belong in the business logic layer?  
How about when a client changes their query without notice? How do you handle neighbor effects?  
This problem doesn't exist in REST because you would do a different endpoint. We ended up putting a REST like cover over GraphQL in APQs.  
Movement to APQ has really just made SIRO "a set of configurable on demand REST endpoints".



## ⚡ REST

- Mature and a known variable
- Quick
- Established in the WM ecosystem
- Versioning
- “Thinner”

## ⚡ GraphQL

- Flexible
- No over/under fetching
- Versioning
- Rapid Iterations

When would I use REST over graphql? Something one off, something that needs to be done quickly and will not necessarily iterate much.

In the general walmart ecosystem suggest evaluating using apollo and just biting the bullet on JS

Other complains, like REST endpoints that to really implement REST have analogs in GraphQL - this may be mitigated with apollo but Java has a long way to go